DATA 180 Recap

Amanda Tran & Chloe Ho

Topics to be covered:

VISUALIZATIONS

DATA WRANGLING

UNSUPERVISED LEARNING

SUPERVISED LEARNING

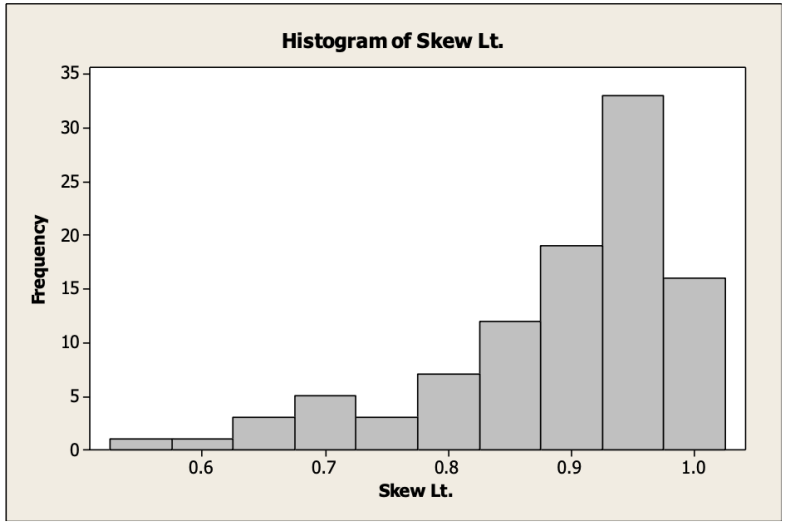

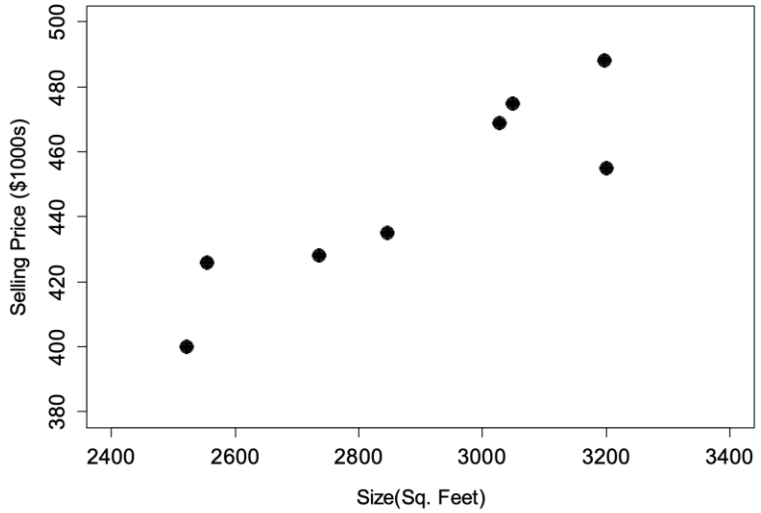

1. Visualizations – Numerical Variables

Single Numerical Variable: Histogram

- Mean, median, mode
- Overall shape of the data set (e.g., symmetric or skewed)
- Presence of (1) gaps in the data set, (2) outliers

Paired data visualization: Scatterplot

1. Visualizations – Categorical Variables

- **Key word: frequency/relative frequency**
- **Visualization examples:** Bar plots & pie charts
- Alternatives… for when there are too many categories? (i.e., balloon plots, mosaic plots)

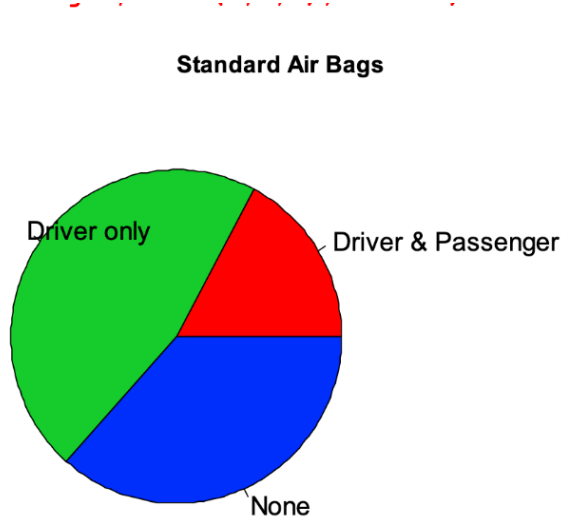

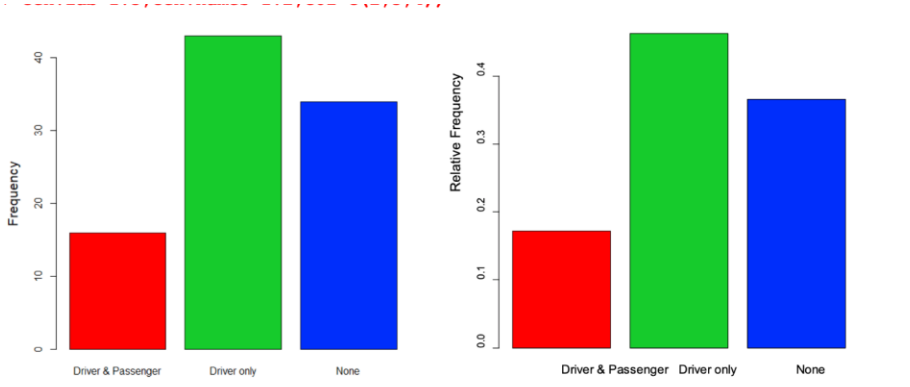

2. Data Wrangling

- **Key Data Wrangling Techniques:**
  - Data Importing: Using functions like read.csv(), read.table()
  - Data Cleaning: Handling missing values, outliers, and incorrect data types
  - Data Transformation: Reshaping data, combining datasets, creating new variables
  - Data Exporting: Saving the cleaned and transformed data for further analysis

- **Why wrangle data?**

- Data Wrangling Trivia

4. Supervised Learning

Observations are classified into ***predictor*** and **response** variables

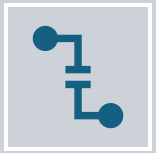

Primary goal: modelling the relationship between a set of predictors and a response variable.

Covered material:

Linear regression

Logistic regression

Simple Linear Regression

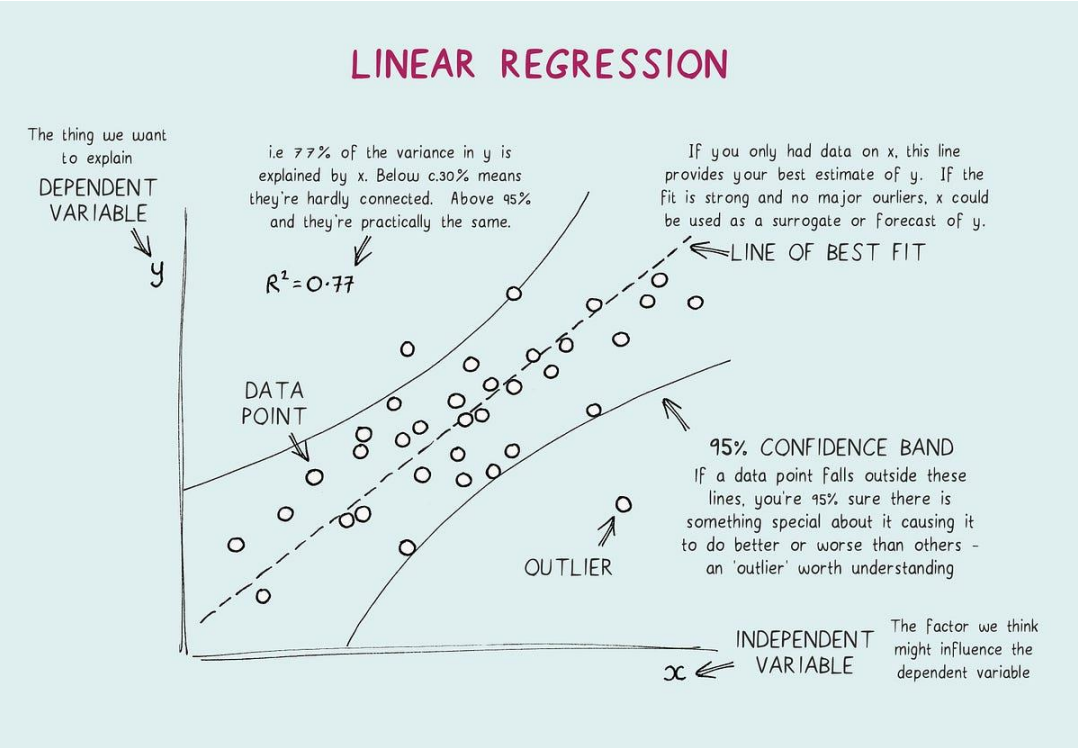

- SLR is a way for predicting a response Y on the basis of a single predictor variable X.
  - Y could be numeric, binary/categorical (coded as integers)
- Easily expandable to Multiple Linear Regression where there exist multiple predictors Xi.
- Eg. $\hat{y} = b_0 + b_1 x$

Logistic Regression

- Logistic regression estimates the probability of an event occurring, such as "voted" or "didn't vote", based on a given dataset of independent variables.
- Since the outcome is a probability, the dependent variable is bounded between 0 and 1.
- *What's the fundamental difference between logistic regression and linear regression?*

3. Unsupervised Learning

For every observation x_i that we observe, we do not observe a response variable y_i.

Goal: finding patterns in the data set.

Covered technique:

K-means clustering

Hierarchical clustering

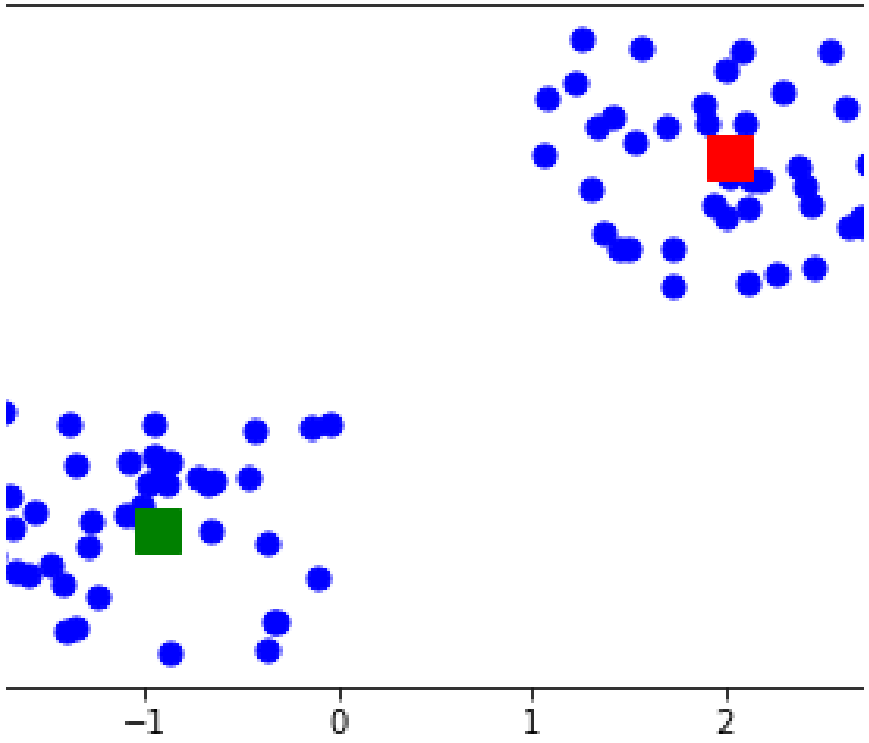

K-means clustering

- You'll define a target number k, which refers to the number of centroids you need in the dataset. A centroid is the imaginary or real location representing the center of the cluster.
- In other words, the K-means algorithm identifies *k* number of centroids, and then allocates every data point to the nearest cluster, while keeping the centroids as small as possible.

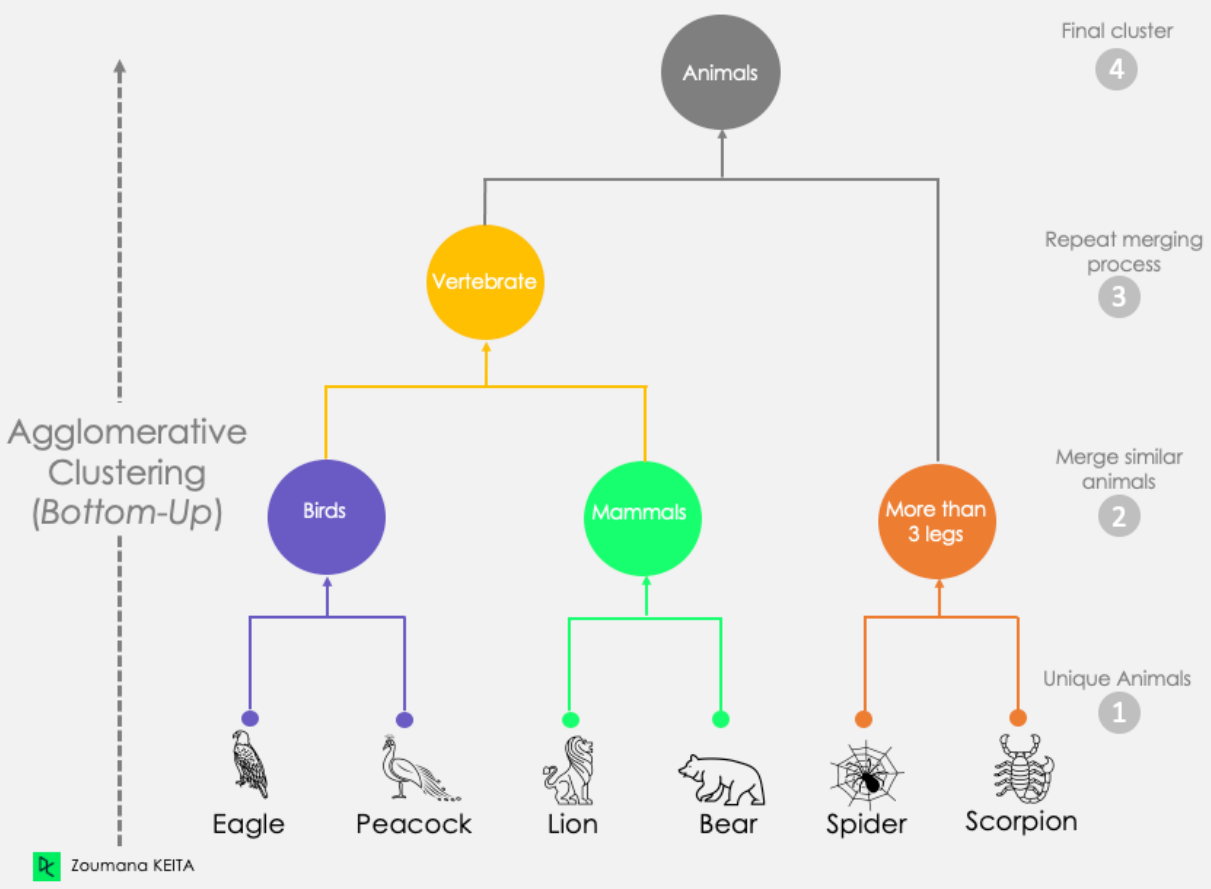

Hierarchical clustering

The basic algorithm of Agglomerative is straight forward:

- Compute the proximity matrix
- Let each data point be a cluster
- Repeat: Merge the two closest clusters and update the proximity matrix
- Until only a single cluster remains

Hierarchical clustering

How do we calculate the similarities between two clusters?

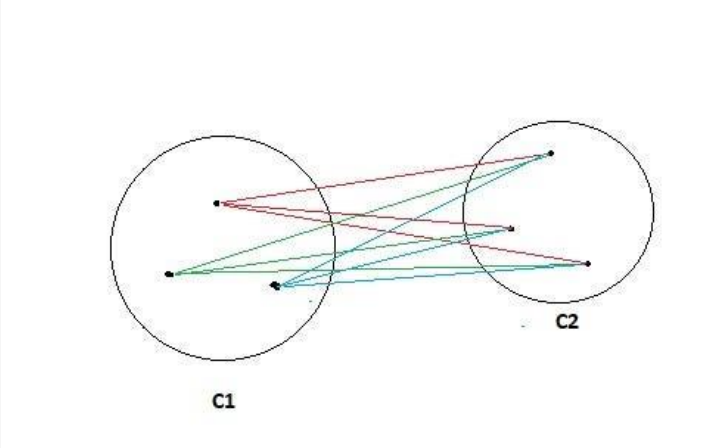

MIN (Single Linkage)

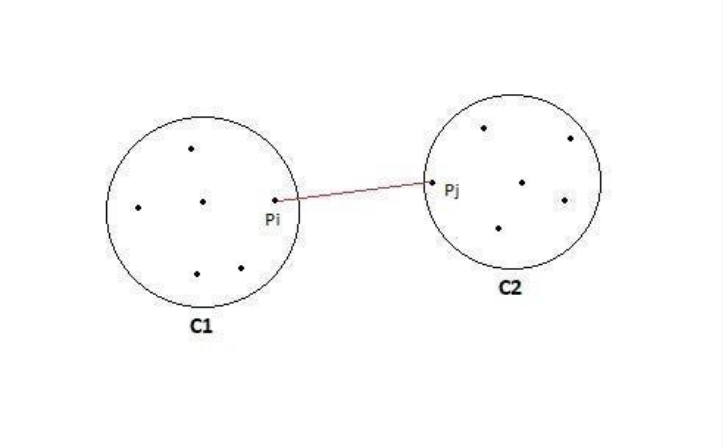

MAX (Complete Linkage)

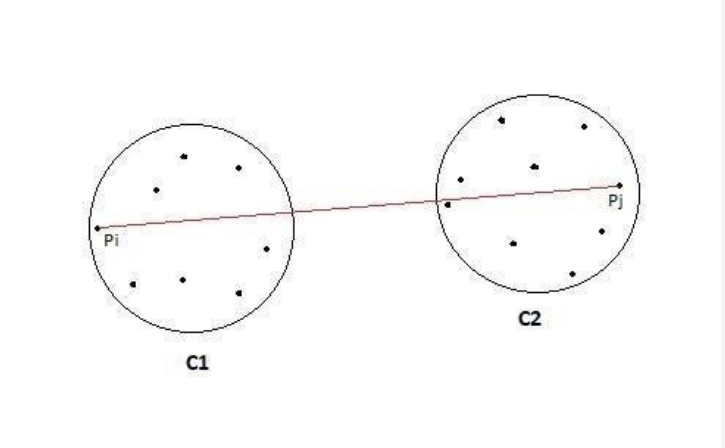

GROUP AVERAGE (Centroid Linkage)